

Machine Learning 1


Hong (PID: A16558957)

2/8/2022

First up kmeans()

Demo od using kmeans() function in base R. First make up some data with a known structure.

```
# Make up some data  
tmp <- c( rnorm(30, mean = 3), rnorm(30, mean = -3) )  
x <- cbind(tmp, rev(tmp))  
  
plot(x)
```



The kmeans() function does k-means clustering

```
k <- kmeans(x, centers = 4, nstart = 20)
k
```

```
## K-means clustering with 4 clusters of sizes 21, 9, 9, 21
##
## Cluster means:
##      tmp
## 1 -3.140785  2.603575
## 2 -1.700734  3.920387
## 3  3.920387 -1.700734
## 4  2.603575 -3.140785
##
## Clustering vector:
## [1] 4 4 3 4 4 4 3 4 4 3 3 4 4 3 4 3 4 4 4 4 3 4 4 4 4 4 3 4 4 4 4 3 4 1 2 1 1 1 1 1 2
## [39] 1 1 1 1 2 1 2 1 2 1 1 2 2 1 1 2 1 1 1 2 1 1
##
## Within cluster sum of squares by cluster:
## [1] 30.05898 11.14108 11.14108 30.05898
## (between_SS / total_SS =  92.6 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

We can use the dollar syntax to get at the results (components of the returned list) > Q1. How many points are in each cluster?

```
k$size
```

```
## [1] 21  9  9 21
```

Q2. What 'component' of your result object details - cluster size? - cluster assignment/membership? - cluster center?

```
k$size
```

```
## [1] 21  9  9 21
```

```
k$centers
```


```
##      tmp
## 1 -3.140785  2.603575
## 2 -1.700734  3.920387
## 3  3.920387 -1.700734
## 4  2.603575 -3.140785
```

```
k$cluster
```

```
## [1] 4 4 3 4 4 4 3 4 4 3 3 4 4 3 4 3 4 4 4 4 3 4 4 4 4 4 3 4 4 4 4 3 4 1 2 1 1 1 1 1 2
## [39] 1 1 1 1 2 1 2 1 2 1 1 2 2 1 1 2 1 1 1 2 1 1
```

Q3. Plot x colored by kmeans cluster assignment and add cluster centers as blue points

```
plot(x , col = k$cluster)
points(k$centers, col = "blue", pch = 15, cex = 2)
```



Hierarchical Clustering


The `hclust()` function needs a distance matrix as input not our original data. For this we use the `dist()` function.

```
hc <- hclust( dist(x) )
hc
```

```
##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

```
plot(hc)
abline( h = 10, col = "red")
```

Cluster Dendrogram



To get our cluster membership vector we need to cut our tree and for this we use the `cutree()`

```
cutree(hc, h = 10)
```

[illegible]

You can cut by a given height $h =$ or into a give number of k groups with $k =$

```
cutree(hc, k=2)
```

[illegible]

Principal Component Analysis

PCA of UK food data

Let's read our data about the weird stuff folkds from the uK eat and drink:

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
```

Look at the first bit of the file:

```
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese           105   103      103       66
## Carcass_meat     245   227      242      267
## Other_meat       685   803      750      586
## Fish            147   160      122       93
## Fats_and_oils    193   235      184      209
## Sugars           156   175      147      139
```

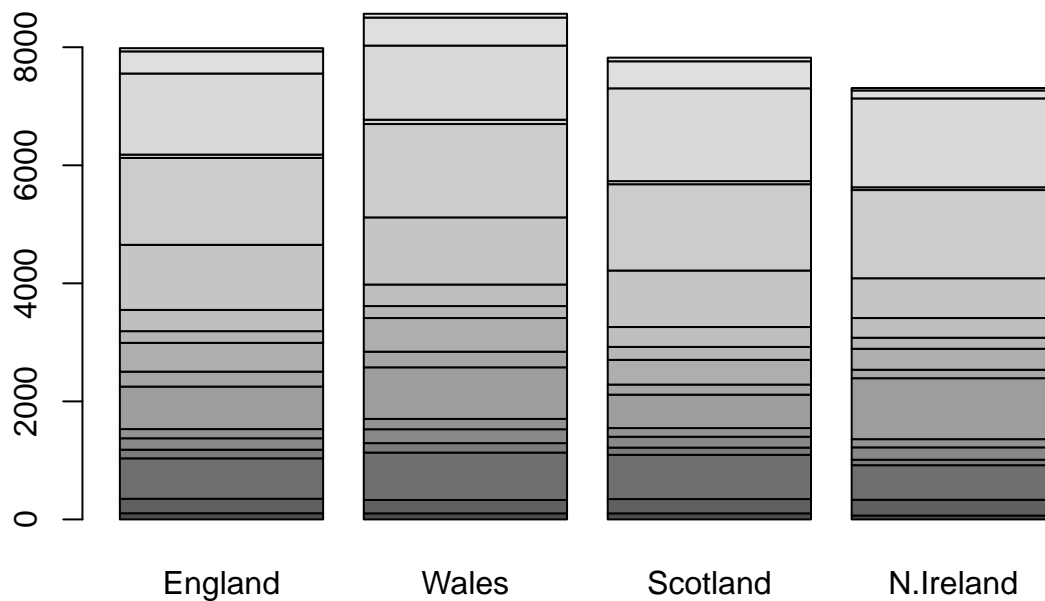
How many columns in this dataset:

```
ncol(x)
```

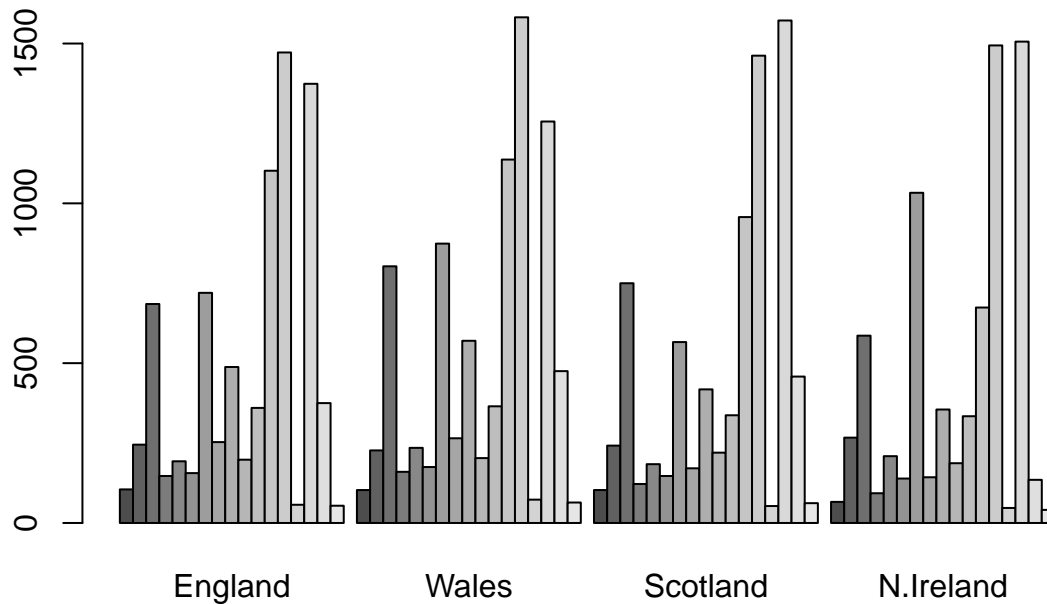
```
## [1] 4
```

We can make some plots to try to understand this data a bit more. For example barplots:

```
barplot(as.matrix(x))
```



```
barplot(as.matrix(x), beside = TRUE)
```



PCA to the rescue

The main base R function for PCA is called `prcomp()`. `prcomp()` expects the observations to be rows and the variables to be columns therefore we need to first transpose our data.frame matrix with the `t()` transpose function.

```
pca <- prcomp( t(x))
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4
## Standard deviation  324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance  0.6744  0.2905  0.03503 0.000e+00
## Cumulative Proportion  0.6744  0.9650  1.00000 1.000e+00
```

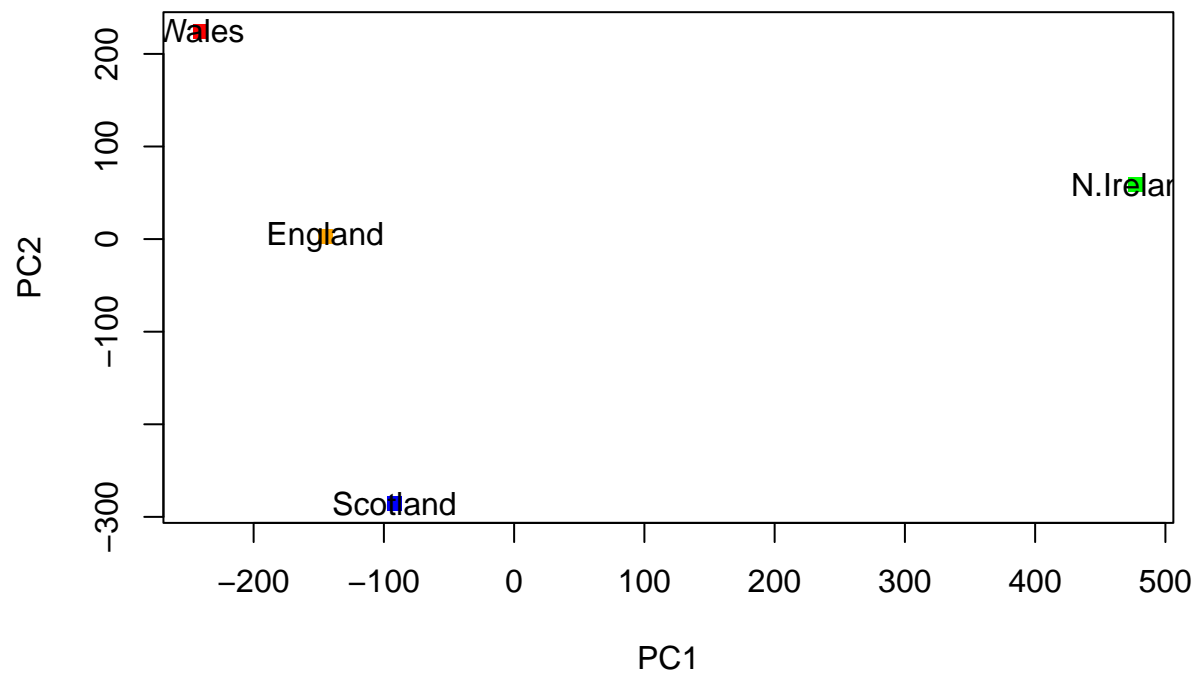
What is in this returned pca object?

```
attributes(pca)
```

```
## $names
```

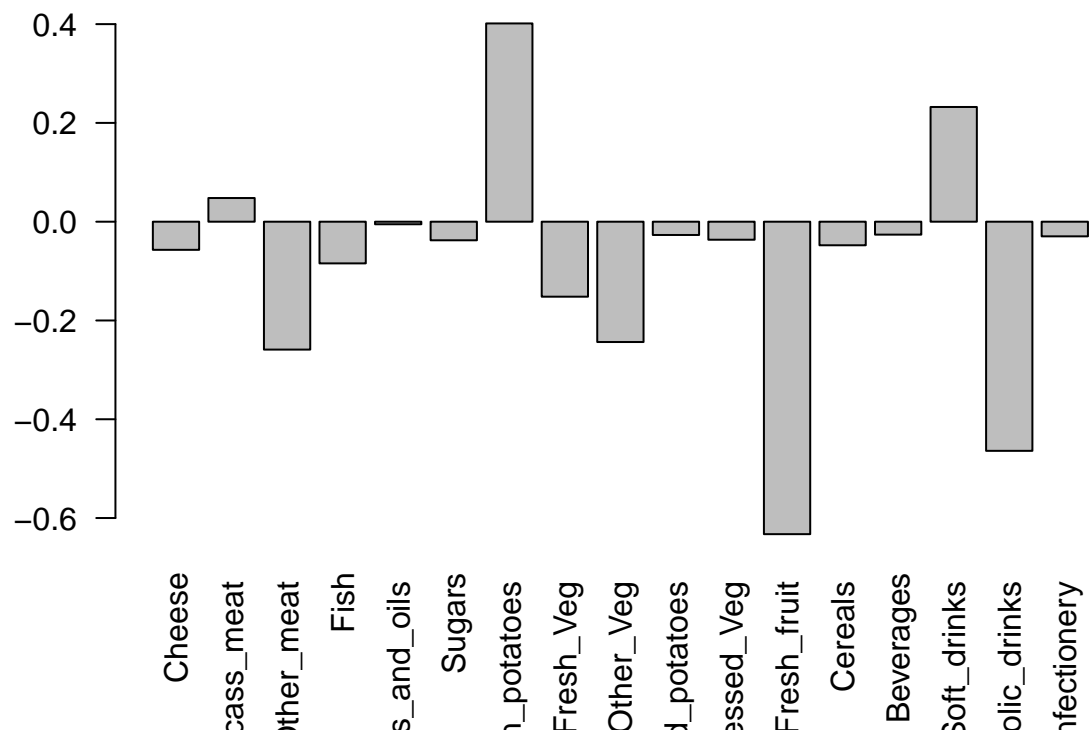
```
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
## $class
## [1] "prcomp"
```

```
plot( pca$x[,1:2] , col = c("orange", "red", "blue", "green"), pch = 15)
text( pca$x[,1], pca$x[,2], labels = colnames(x))
```



We can look at how the variables contribute to our new PCs by examining the `pca$rotation` component of our results.

```
barplot(pca$rotation[,1], las = 2)
```



An RNA-Seq PCA example...

Read the data first

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##          wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1    439 458  408  429 420  90  88  86  90  93
## gene2    219 200  204  210 187 427 423 434 433 426
## gene3   1006 989 1030 1017 973 252 237 238 226 210
## gene4    783 792  829  856 760 849 856 835 885 894
## gene5    181 249  204  244 225 277 305 272 270 279
## gene6    460 502  491  491 493 612 594 577 618 638
```

How many genes (how many rows)?

```
nrow(rna.data)
```

```
## [1] 100
```

How many experiments (columns)?


```
ncol(rna.data)
```

```
## [1] 10
```

Let's do PCA of this dataset. First take the transpose as that is what the `prcomp()` function wants...

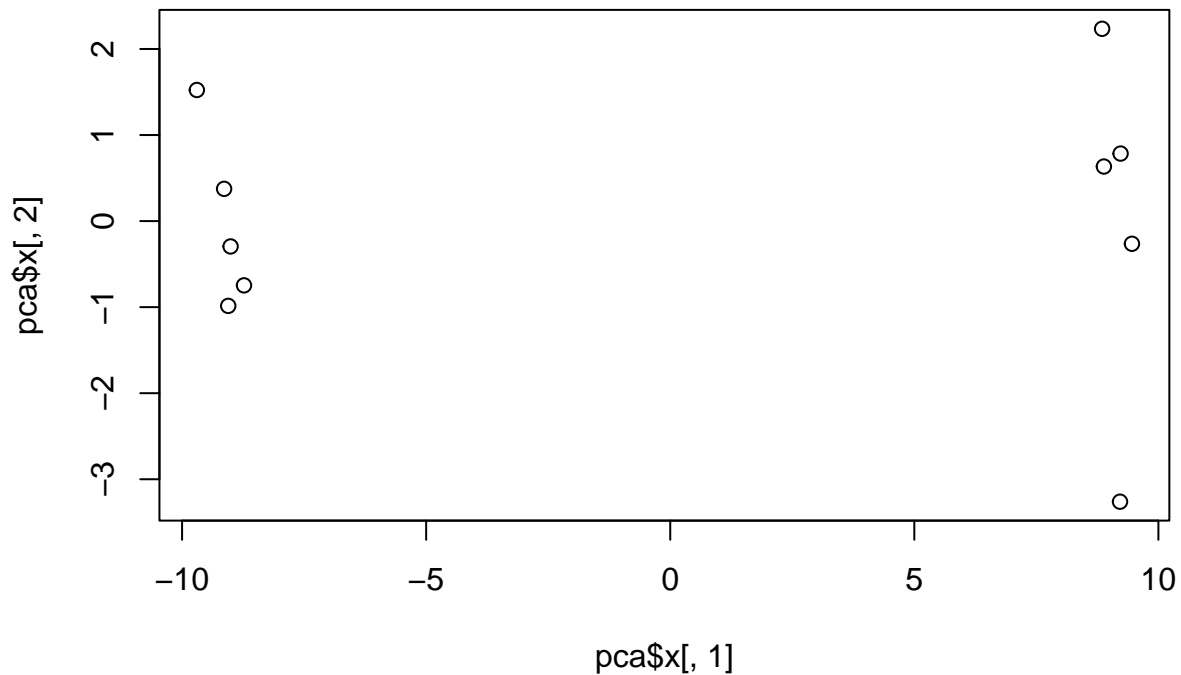
```
pca <- prcomp( t(rna.data), scale = TRUE)
summary(pca)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237  1.5198  1.05787  1.05203  0.88062  0.82545  0.80111
## Proportion of Variance 0.9262  0.0231  0.01119  0.01107  0.00775  0.00681  0.00642
## Cumulative Proportion 0.9262  0.9493  0.96045  0.97152  0.97928  0.98609  0.99251
##              PC8      PC9      PC10
## Standard deviation  0.62065  0.60342  3.348e-15
## Proportion of Variance 0.00385  0.00364  0.000e+00
## Cumulative Proportion 0.99636  1.00000  1.000e+00
```

We can make our score plot (a.k.a. PCA plot) from the `pca$x`

```
plot(pca$x[,1], pca$x[,2])
```



Make a little color vector to color up our plot by `wt` and `ko`

```
colvec <- c(rep("red",5), rep("blue",5))  
plot(pca$x[,1], pca$x[,2], col = colvec, pch = 15)
```

